Dissertation Review — marker's assessment and required corrections

Prepared 3 September 2026. Every claim below was checked against the code and the analysis outputs, not against the project's own documentation.

1. Where this sits

HONEST VERDICT — as a Sheffield MSc AI dissertation marker.

WHERE THIS SITS. On the written outputs as they stand, this is a solid upper-second project with a clear route to a distinction that it has not yet taken. The engineering and the framing are distinction-level. The statistics are not, and the gap between how rigorous the documents SAY they are and how rigorous they actually are is the single biggest risk — because the documents draw attention to their own rigour, which invites the examiner to check it, and it does not survive checking.

STRENGTHS (real, not consolation). 1. The question is genuinely novel and genuinely narrow. "To what extent does an automatic image-emotion reader agree with prompt-derived emotion labels for AI-generated monochrome drawings" is a well-formed, answerable MSc question, and the dossier's paragraph explaining why it is deliberately narrower than "does the drawing convey the feeling" is the best paragraph in the project. 2. The dataset is not downloadable, the prompts are real first-person self-report, and the engineering to make it usable — three filename schemes, C2PA timestamp recovery identifying the 13 pre-festival tests — is exactly the kind of unglamorous work that distinguishes a real project from a Kaggle exercise. Recovering developer tests from signing timestamps is a nice piece of forensics and should be foregrounded. 3. "Chance is not zero" is the strongest single idea here. Discovering that the headline metric has a floor of 41% and a random-pairing expectation of 76%, and then rebuilding the reporting around that, is a genuine methodological insight, honestly arrived at, and it will earn credit in the viva. 4. The rejected-approaches table is evidence, not excuse-making. Qwen thrashing 16 GB with zero progress in 4 h 45 min, Florence-2's transformers incompatibility, GoEmotions collapsing on third-person descriptions, a bare caption collapsing to "surprise" — these are findings and they belong in the methods chapter exactly as written. 5. The withdrawal of the positive-versus-negative claim, on the grounds of a degenerate null and two readers biased in opposite directions, is the mark of someone doing research rather than producing a result. Say so out loud in the viva; it is worth more than the finding would have been. 6. Name redaction: reviewed decisions with recorded reasons, the "Jack Frost" over-reach, and the dedup-key fix that stopped two different people's prompts merging into one. That last one is a subtle bug caught properly. 7. Reproducibility infrastructure — fixed seeds, resumable JSONL, a runnable self-test, a live demonstrator — is well above the norm at this level.

WEAKNESSES. 1. There is no inferential statistics in this project. I grepped: no scipy, no bootstrap, no chi-square, no Fisher, no confidence interval, no paired test anywhere in the codebase. Every quantity is a point estimate with a permutation z against its own null. The headline finding is a DIFFERENCE between two groups, and the difference is never tested. "The gap survives every check" names checks that do not exist. The dossier's limitation 6 admits this, which is honest, but admitting a missing analysis does not substitute for running it — and a bootstrap CI on the face-scene difference is roughly forty lines of code and an afternoon. 2. The negative control, which the dossier explicitly nominates as mattering "more than any single result", does not do what it claims. Two of its four metrics are degenerate by construction (every prompt in the subset shares the neutral label, so permutation cannot move family agreement — I reproduced z = +0.0 exactly), and the subset is contaminated: at minimum 15% of it contains an explicit emotion word ("ecstatic", "delighted but tired", "content", "hopeful", "inspired"), nearer 40% by inspection, and two records are the bare token "[name]". The control is not worthless — the VA result at z=+0.3 is meaningful — but it is one metric on a noisy subset, not four metrics on a clean one. 3. The circularity caveat is wrong about its own mechanism. This is the most damaging finding in my review, precisely because it is the caveat. The write-up says scene descriptions "lack expression content"; 96.4% of them contain a readable expression word. The split is driven by the people-count answer, and the actual artifact is that the description template writes "A drawing of none showing…", after which the NLI reader returns neutral for 35% of emotional scenes versus 5% of faces — and a neutral image reading cannot match a signed prompt. Fixing this is straightforward and would materially strengthen the chapter; leaving it will not survive a competent second marker. 4. Two headline findings rest on unvalidated classifier splits. Neither the emotional/non-emotional split nor the face/scene split has a precision or recall figure, despite PIPELINE_DESIGN.md having planned exactly that ("validate on ~50-100 hand-labelled

images"). Hand-labelling 100 images and 100 prompts is a day's work and would convert both splits from assumptions into measurements. 5. The consistency finding is the only headline with no null, in a project whose central methodological claim is that every number is quoted against a null. I computed it: the null is 59.4%, so 68.8% is z~+4.4 — still a finding, but a quarter as dramatic as an unbaselined "69%" sounds, and half the groups are pairs with a 50% floor. Worse, the variance is attributed entirely to DALL-E when it also contains BLIP-VQA and NLI noise for which no reliability estimate exists. 6. The architectural claim is softer than presented. "The vision model is never asked for an emotion" is literally true of the question wording and substantively misleading: the expression question returns a bare emotion word for 38% of records, and the "separate" text model reproduces that word's family 69% of the time. The method comparison may therefore be showing that BLIP-VQA reads line art better than CLIP, not that describe-then-extract is a better architecture — and the comparison additionally uses a different prompt reader and a different subset definition from every other headline, which neither document discloses. 7. Cross-document drift. Fourteen quantities are stated differently across the six files, including a factor-of-ten-in-credibility one (66% versus 57% for the same statistic, with the inflated version in the file labelled authoritative), and a claim that the test suite has 28 checks when it has 29 smoke checks that assert no numbers at all.

WHAT WOULD SPECIFICALLY LOSE MARKS. - Any perception verb surviving into the final text. "Convey" and "carry" in the README, and above all in the three figure titles, contradict the dissertation's own most careful paragraph. A marker who reads the figure titles and then the limitations chapter will conclude the student does not believe their own caveats. This is the cheapest fix in the whole list and the most expensive to leave. - "Detected automatically and no chance figure is reported for them." It is a hand-passed boolean on two of six code paths, and the z is printed anyway. A false claim about a safeguard is worse than no safeguard, because it converts a limitation into an integrity question. - "All survive, at z between +9 and +11", contradicted by the same document four pages later. This will read as selective reporting even though the correct number is stated elsewhere. - "Independently reviewed for statistical validity" with no reviewer, date or scope. This is the first thing an examiner will ask about, and there is no good answer. - Absent confidence intervals on every reported proportion. At MSc level this is now expected, not optional, and the omission is visible on every table. - The face/scene headline framed as a claim about faces rather than about a classifier's output. The dossier's own "Stated carefully" paragraph gets this right; the heading above it, the README, the presentation and the figure title all get it wrong. Markers weight headings. - "Twelve themes" when the bibliography has eight, and "28 tests" when there are 29. Small, but they are the errors a marker can verify in one minute, and they calibrate how much of the rest they will trust.

WHAT TO DO WITH THE REMAINING TIME, IN ORDER. (1) Bootstrap CIs on every reported proportion and a bootstrap or chi-square test of the face-scene difference — one afternoon, and it removes the single largest weakness. (2) Rewrite the circularity caveat to describe the actual mechanism (people-count clause, "A drawing of none" prefix, neutral-cap on valence sign) — one hour, and it turns a wrong caveat into a genuinely sophisticated one. (3) Compute and report the null for the consistency finding — twenty minutes. (4) Hand-label 100 prompts and 100 images to give both splits a precision/recall figure — one day, converts two assumptions into measurements. (5) Do a global find-and-replace on the perception verbs and the terminology list above — two hours. Items 1-3 alone move this from a good 2:1 to a credible distinction, because they close the gap between claimed and actual rigour, and that gap — not the absence of a human study, which is properly justified — is what is currently holding the work back.

---

2. Claims that must be corrected before they enter the dissertation

These were found to overreach their evidence. Each is given with the correction to use instead.

1. "does the drawing convey the feeling in the words?" and "Faces carry the emotion; scenes carry it far more weakly"

Where: emoviz/README.md (opening paragraph, and Key results #1); emoviz/outputs/RESULTS_final.md (Headline finding 1 heading) **Problem:** "Convey" and "carry" are perception verbs — they assert something about a viewer. No viewer exists in this project. The measured quantity is the rate at which a BART-MNLI reading of a BLIP-VQA description lands in the same Ekman family as a GoEmotions reading of the prompt

text. The dossier states this correctly ('It does not establish that a human viewer perceives the same feeling'), so the README and RESULTS_final contradict the project's own stated scope rather than merely being loose. **Use instead:** "Does an automatic image-emotion reader agree with a prompt-emotion reader?" / "Prompt-image label agreement is substantially higher for drawings the describer classifies as containing a person with a readable expression than for those it does not."

2. Figure titles: "Drawings with a readable face convey the prompt's emotion far better than scenes"; "The describe->extract method reads image emotion fa

Where: emoviz/make_figures.py lines 63-64, 84-85, 103-104 — rendered into figs 1-3 and reproduced in PROJECT_EXPLAINED.pdf / project_dossier.tex ('Findings', 'A negative control that passes') **Problem:** These three figure titles are the most over-claiming sentences in the project, and they sit inside the document whose body text carefully disclaims exactly those readings. Markers read figure titles. 'Convey' is perception; 'reads image emotion far better' presupposes a ground truth for image emotion that does not exist (only agreement with a second model was measured); 'appears only where there is a feeling to convey' is false, per the contaminated control below. **Use instead:** Fig 1: "Prompt-image label agreement is higher for face-classified than scene-classified drawings (n=402 / 205)". Fig 2: "Describe-then-extract agrees with the prompt reader more often than direct CLIP (same 608 records, NLI prompt side)". Fig 3: "Agreement is indistinguishable from chance on prompts the text reader labels neutral".

3. "scene descriptions lack expression content, which mechanically lowers their scores"; "the face/scene split is defined by whether BLIP-VQA reported an

Where: emoviz/outputs/RESULTS_final.md (Headline 1, 'Honest caveat — partial circularity'); project_dossier.tex Limitation 2; PRESENTATION_TECHNICAL_STORY.md 'Limits to state yourself' **Problem:** Factually wrong, and it is the project's flagship honesty statement, which makes it doubly damaging. I checked: 406 of the 421 'scene' records (96.4%) DO contain a readable expression word — 'smile' x161, 'smiling' x92, 'happy' x79, 'angry' x31, 'sad' x20. They are classified 'scene' almost entirely because BLIP-VQA answered 'none' or '0' to 'How many people are there?' (414/421). The rule is a conjunction and the people-count clause does essentially all the work. The real mechanism is different: the composed description opens "A drawing of none showing…", and BART-MNLI then returns neutral as top-1 for 35.1% of emotional-scene descriptions against 5.2% for faces. Because valence-sign agreement scores a neutral image reading as a miss, the scene rate is mechanically capped near 65%. The caveat therefore misstates both the rule and the artifact. **Use instead:** "The split is not a face detector. It is driven almost entirely by BLIP-VQA's people-count answer: 414 of the 421 scene records are ones where it answered 'none' or '0', while 96% of them still carry an expression word. The description for those records opens 'A drawing of none…', and the NLI reader returns neutral for 35% of them against 5% for face records. Since a neutral image reading cannot match a signed prompt, a large share of the face/scene gap is an artifact of the description template and the metric's treatment of neutral, not evidence about faces."

4. "faces align roughly 1.7x better on valence and family"; "Face drawings align about 1.7 times better on both sign and family"; "gap survives every che

Where: RESULTS_final.md Headline 1; project_dossier.tex 'Alignment is real, and much stronger for drawings containing a face'; PRESENTATION_TECHNICAL_STORY.md finding 1 **Problem:** Three problems. (a) There is no statistical test of the face-vs-scene difference anywhere in the codebase — I grepped: no scipy, no bootstrap, no chi-square, no Fisher, no confidence interval, no paired test. Every z reported tests 'this subset beats its own null', never 'these two subsets differ'. So 'the gap survives every check' names checks that were never run on the gap. (b) The 1.7x ratio uses raw rates while the two subsets have very different nulls (face null 53.5%, scene null 29.6%); measured as excess over chance the gap is 20.4pp vs 13.3pp, i.e. ~1.5x, and on the VA metric the two are indistinguishable. (c) With four metrics across six subsets across two samples, no multiple-comparison control is applied, yet Figure 3 draws a 'significance threshold (z~2)' line. **Use instead:** "Face-classified drawings agree with the prompt reader on 73.9% of valence signs against a chance level of 53.5%; scene-classified drawings on 42.9% against 29.6%. Both exceed their own nulls (z+10.4 and z+5.8). The difference between the two groups has not itself been tested — a chi-square or bootstrap on the difference,

with a confidence interval, is required before the gap is claimed, and the ratio of raw rates overstates it because the two nulls differ."

5. "405 of the 1,012 prompts contain no feeling at all — 'pigs', 'mickey mouse', 'disney characters' "; "~400 prompts contain no feeling … It only finds e

Where: project_dossier.tex 'A negative control that passes'; PRESENTATION_TECHNICAL_STORY.md section 6; RESULTS_final.md Headline 2 **Problem:** The control set is defined as 'GoEmotions top-1 = neutral' — a classifier output, never validated against human labels (PIPELINE_DESIGN.md proposed hand-labelling ~50 prompts for precision/recall; no such validation is reported). It is visibly contaminated. By a conservative emotion lexicon, 62 of 405 (15.3%) contain an explicit emotion word: 'ecstatic', 'delighted but tired', 'content' x2, 'calm' x2, 'Hopeful', 'inspired', 'hungover' x3, 'tired, hungry', 'relaxed but with menopausal hot flashes' x4, 'nostalgic over a reality you never got to experience'. In a random sample of 40 the true rate of affect-bearing prompts is nearer 40% once 'jealous', 'boring', 'chipper', 'bonkers', 'gnarly', 'angry and comfused', 'demure and mindful' and a French prompt about stress and joy are counted. Two records in the control are literally the string '[name]' (records 750, 751 — originally two different people's full names); redaction consumed the whole prompt and they were left in. The three examples quoted in the dossier are the cleanest cases in the set. **Use instead:** "405 prompts are labelled non-emotional by the text reader's top-1 output; on these the pipeline shows no above-chance agreement. Two caveats limit what this establishes: the label is the text reader's own judgement, never validated against human coding — at least 15% of the set contains an explicit emotion word ('ecstatic', 'content', 'hopeful') — so the subset is better described as 'prompts the text reader could not read affect from' than as 'prompts with no feeling'; and two records reduce to '[name]' after redaction and should be dropped."

6. "emotion-family agreement sits exactly on its null" — presented under the heading "Why this matters more than any single result"

Where: project_dossier.tex 'A negative control that passes' **Problem:** It sits exactly on its null because it cannot do otherwise. Every prompt in that subset has top-1 = neutral by the subset's own definition, so family agreement is simply the proportion of image readings that are neutral — a quantity permutation of the image side leaves numerically unchanged. I reproduced it: true 0.286, null 0.286, z = +0.0 exactly; valence sign 0.000 vs 0.000. This is an algebraic identity reported as an empirical validity result. RESULTS_final.md spots this for the valence-sign cell ('a definitional artifact, not a result') but not for family, and the dossier presents both as passing checks. Only VA carries information here, and it gives z=+0.3. **Use instead:** "On the non-emotional subset only the valence-arousal metric is informative: 76.6% against a null of 76.6% (z=+0.3). Valence-sign and family agreement are degenerate here — every prompt shares the neutral label, so permutation cannot change them and their z is identically zero by construction, not by measurement. The control therefore rests on one metric, not four."

7. "Degenerate nulls are refused. Where a subset shares one label by construction … These are detected automatically and no chance figure is reported for

Where: project_dossier.tex 'How the numbers are kept honest' §'Degenerate nulls are refused' **Problem:** Both halves are false and checkable in twenty seconds. In emoviz/analysis_final.py, `degenerate` is a hand-passed boolean argument to `sub_analysis`, supplied manually only for the positive- and negative-prompt blocks — nothing is detected. And when it is passed the z is still printed; the function merely appends '(single-sign subset: sign/fam nulls DEGENERATE, ignore their z)'. The non-emotional subset, equally degenerate, is not flagged at all and its z-scores are reported as a headline. Claiming an automatic safeguard that is a manual annotation on two of six code paths is precisely what an examiner checks. **Use instead:** "Two subsets are degenerate for the sign and family metrics because every prompt in them shares one label: the single-valence subsets and the non-emotional subset. In these cases permutation cannot change the score and the z is uninformative. These are currently identified by hand in the analysis script and annotated in the output rather than suppressed; automatic detection — refusing to print a z when the prompt side has one distinct top-1 label — is an outstanding fix."

8. "Every headline was re-run on the 763 distinct prompts only. All survive, at z between +9 and +11."

Where: project_dossier.tex 'A robustness check on repeated prompts' **Problem:** Contradicted twice within the project's own documents and by the data. The dossier's own Findings section reports distinct-prompt scene agreement at z+3.5, and RESULTS_final.md says plainly 'scene weakens (z+3) and negative drops to non-significant'. My replication on the 763 distinct prompts gives face sign z+9.0 but scene sign z+3.4 and scene family z+2.4 — the last marginal at any conventional threshold and unlikely to survive correction for the number of comparisons. Writing 'all survive at z between +9 and +11' in the rigour section, four pages before reporting z+3.5 for one of those headlines, reads as carelessness at best and selective reporting at worst. **Use instead:** "Every headline was re-run on the 763 distinct prompts. Overall agreement and the face result are essentially unchanged (z +9.3 to +10.8; face valence sign z +8.4). The scene result weakens substantially (valence sign z +3.5, family agreement z +2.6, n falling from 205 to 133) and should be treated as suggestive rather than established. The single-valence subsets have degenerate nulls and are not re-testable."

9. "Generative emotion is unstable: given the same prompt, the model conveys a different feeling roughly a third of the time"; "The generator is inconsis

Where: RESULTS_final.md 'Extra finding — generative consistency'; emoviz/README.md Key result 3; project_dossier.tex 'The generator is inconsistent with itself'; PRESENTATION_TECHNICAL_STORY.md finding 3 **Problem:** Three defects, and this is the finding most exposed in a viva. (a) The 'emotion' compared across repeats is `image_top1` — the pipeline's own reading. The measured 31% disagreement is generator variance PLUS BLIP-VQA variance PLUS NLI variance, and no reliability estimate exists for the reader (no test-retest, no second describer, no human check). The variance is attributed entirely to DALL-E without warrant. (b) It is the only headline reported with no null, in direct contradiction of the dossier's own claim that 'every number reported in this project is quoted against its own null'. I computed it: the permutation null for mean modal agreement is 59.4% (sd 2.1), so 68.8% is z ~ +4.4 — real but far less dramatic than 69% sounds; the null for fully-consistent runs is 14.5, so 28/87 is z ~ +4.0, not the damning number it is presented as. (c) The 59% floor is high because 43 of the 87 groups are pairs, whose modal agreement cannot fall below 50%. Quoting 69% without that floor repeats exactly the error the project is praised for catching in its own section 7, 'Chance is not zero'. **Use instead:** "Across 87 prompts drawn two or more times, the pipeline reads the same emotion for a mean 68.8% of drawings within a group, against a permutation null of 59.4% (sd 2.1; 10,000 draws, seed 7; z +4.5); 28 of 87 groups are fully consistent against a null of 14.5 (z +4.2). Half the groups are pairs, for which modal agreement has a floor of 50%. This measures the stability of the whole prompt->image->reader chain, not of the generator alone: with no reliability estimate for the reader, an unknown share of the variation is measurement noise."

10. "given the same prompt, does DALL-E draw the same feeling?"

Where: RESULTS_final.md 'Extra finding'; emoviz/data_audit_consistency.py RQ3 print block **Problem:** The groups pool text typed by different visitors on different days with consecutive re-generations by one visitor. Sampling temperature, seed, model version and any system-prompt scaffolding used by the festival installation are unknown and unlogged. This is an observational grouping over a black box, not a controlled re-generation experiment. The documents do disclose the pooling — creditable — but the research question stated on top of it exceeds what the pooling permits. **Use instead:** "Where the same prompt text appears more than once in the corpus — whether typed by different visitors or re-submitted by one — how stable is the emotion the pipeline reads from the resulting drawings? Generation parameters were not recorded, so this bounds the reproducibility of the deployed system as a whole rather than isolating the generator."

11. "The vision model is never asked for an emotion"; "emotion is inferred from its factual description by a separate text model"; "The word 'emotion' nev

Where: emoviz/README.md ('Describe, don't ask'); PRESENTATION_TECHNICAL_STORY.md section 4; project_dossier.tex 'Stage by stage' step 2 **Problem:** The literal claim (the token 'emotion' is absent from the five questions) is true; the claim it supports is not. 'What is the facial expression?' elicits a bare emotion word in 38% of records: happy x207, angry x75, sad x71, surprised x25, excited x3, confused x3. And the downstream 'separate text model' largely copies it: where the expression answer maps to a single emotion,

BART-MNLI's top-1 reproduces that emotion 68.7% of the time. The architectural separation is therefore substantially cosmetic — which matters, because the headline method comparison is presented as vindicating that architecture. The honest reading is that BLIP-VQA is a better affective reader of line art than CLIP, a finding about models rather than about architecture. **Use instead:** "The vision model is asked only factual questions and the string 'emotion' appears in none of them. This is a weaker separation than it sounds: the facial-expression question returns a bare emotion word ('happy', 'angry', 'sad') for 38% of records, and where it does, the downstream text model reproduces that word's family 69% of the time. The advantage over direct CLIP may therefore lie in BLIP-VQA being a better affective reader of line art rather than in the describe-then-extract architecture; an ablation that strips emotion words from the expression answer before the extraction stage would separate the two."

12. "New method beats the old one"; "describe->extract reaches 49.3% emotion-family agreement against 25.3% for asking CLIP directly … The redirection was

Where: RESULTS_final.md Headline 4; project_dossier.tex 'The method matters, and the supervisor's architecture wins'; README Key result 4 **Problem:** Four issues. (a) 'Beats' and 'wins' presuppose a criterion of correctness; there is none. The new method agrees more with a text model reading the prompt — if the text model is wrong, higher agreement is worse. Without ground truth this is a similarity measurement, not a benchmark victory. (b) similarity_v2.py holds the prompt side constant at `r['prompt_emotions']` — the BART-MNLI prompt reader — not the GoEmotions reader used for every other headline, so 49.3% and the headline 53.1% are not the same measurement, and neither document says so. (c) The 608-record subset is selected by the v1 pipeline's `is_emotional` flag, which comes from a hand-built lexicon in emoviz/text_emotion.py — a different definition of 'emotional prompt' from the GoEmotions rule used elsewhere. (d) The dossier's own limitation 6 admits there is no paired test; two point estimates with no CI and no paired statistic do not support 'wins'. **Use instead:** "Holding the prompt reader constant (BART-MNLI, not the GoEmotions reader used elsewhere) and varying only the image reader on the 608 records both pipelines cover: describe-then-extract agrees with the prompt reader on 49.3% of emotion families against 25.3% for direct CLIP. The subset is selected by the v1 lexicon's definition of an emotional prompt; these figures are not comparable with the 53.1% headline; and no paired test or confidence interval has been computed. Neither reader has ground truth, so this shows the two methods disagree, with describe-then-extract closer to the text reader — not that it is more accurate."

13. "All 1,025 images were processed (100%, 0 failures)"; "1,025 records, 0 failures"; test.sh assertion "no processing failures"

Where: RESULTS_final.md 'Data accounting'; RESULTS_describe_extract.md header; project_dossier.tex dataset table; emoviz/test.sh line 22 **Problem:** '0 failures' means zero unhandled exceptions. It is used throughout as if it were a statement about output quality. BLIP-VQA returned a confident facial-expression answer for all 421 records in which it simultaneously reported zero people present — 332 of them 'smile'/'smiling'/'happy'. That is a ~32% corpus-wide hallucination rate on the single answer the entire pipeline depends on, never quantified, while 'zero failures' is repeated in four documents. There is no validation of any BLIP-VQA answer against human labels, and no precision/recall for either the emotional/non-emotional split or the face/scene split — both of which PIPELINE_DESIGN.md said would be hand-validated on 50-100 items. **Use instead:** "The pipeline ran to completion on all 1,025 images with no unhandled errors. This is a robustness statement, not an accuracy one: the describer returns a facial-expression answer even where it reports no people present (332 of those 421 records answer 'smile', 'smiling' or 'happy'), and none of its answers has been validated against human coding. Precision and recall for both the prompt split and the face/scene split are unmeasured, and are the principal unquantified source of error in the results."

14. "the strongest evidence says lightness and chroma — not hue, and not shape — carry most of the emotional meaning in text-to-image output"

Where: project_dossier.tex 'Why the question is open', second bullet **Problem:** The cited evidence (Alvarado 2025, summarised in RESEARCH PAPERS/RESEARCH_GAP.md and ANNOTATED_BIBLIOGRAPHY.md line 314) extracts CIELAB features — Chroma, L*, a*, b** — and maps them to VAD. It contains no shape variable. 'Not shape' is therefore not a finding of that paper but an inference from a variable's absence — exactly the reasoning error the dossier is elsewhere careful about. And the sentence is load-bearing: 'colour carries the

meaning and shape does not' is the entire motivation for treating monochrome line-art as a hard case. **Use instead:** "Alvarado (2025) finds that in text-to-image output lightness and chroma, rather than hue, are the principal colour carriers of emotional meaning. That study models colour features only; it does not test shape, so it gives no evidence either way about whether form and line can carry affect. That open question is what makes this monochrome corpus interesting."

15. "fear / disgust / surprise are poorly captured by both readers -> restrict strong claims to the positive/negative (valence) axis"; "likely genuinely ha

Where: RESULTS_final.md 'What we CANNOT claim'; RESULTS_describe_extract.md 'Honest limitations' **Problem:** Two problems. The second quote attributes reader failure partly to the absence of colour, in a corpus with no colour condition — precisely the claim the dossier's 'important restraint' paragraph forbids. And 'restrict strong claims to the valence axis' contradicts the bullet immediately above it, which withdraws all positive-vs-negative claims because the single-valence nulls are degenerate and the two readers point in opposite directions. The document simultaneously tells the reader to trust and not to trust the valence axis. Note also that the 'surprise' family is dominated by curiosity-mapped prompts under the standard GoEmotions->Ekman map, so its low recovery is partly a taxonomy artifact. **Use instead:** "Fear, disgust and surprise are recovered at very low rates by both readers, on subsets of n=16 (fear) and n=1 (disgust) under the final GoEmotions prompt reader, and a surprise class largely composed of curiosity-mapped prompts. With no colour condition, none of this can be attributed to the absence of colour; it is equally consistent with class rarity, the GoEmotions->Ekman mapping, and reader bias. Claims are therefore restricted to overall agreement and the face/scene contrast; no directional claim is made on the valence axis, whose single-valence subsets have degenerate nulls."

16. "At the Festival of the Mind 2024, about a thousand members of the public typed a few words about how they were feeling"; "1,025 short public self-rep

Where: project_dossier.tex opening key box; emoviz/README.md opening paragraph **Problem:** 1,025 counts records, not people. Thirteen are developer tests, 249 repeat an earlier prompt, and the project's own consistency section states the repeats include 'consecutive re-generations' by a single visitor. The number of contributing individuals is unknown and bounded above by 763. Separately, 'self-reports of emotion' is inconsistent with the same project's claim that 405 of them 'contain no feeling at all' — the corpus cannot be both. The installation evidently accepted any text ('pigs', 'birds eye view of sheffield from space', 'square a circle with a pair of compasses and a straight edge'), which is a property of the data collection that should be stated. **Use instead:** "The corpus comprises 1,025 prompt-drawing records collected at Festival of the Mind 2024, of which 1,012 are analysed. Visitors were invited to describe how they were feeling but the interface accepted free text, and a substantial minority of entries are object or scene requests rather than self-report. Records are not linked to individuals and 249 repeat an earlier prompt, so the number of contributors is unknown and no greater than 763."

17. "Redaction is analytically inert: re-running the text-emotion model on the redacted prompts moves each probability by at most 0.012 and changes zero t

Where: project_dossier.tex 'Name redaction'; emoviz/README.md 'Name redaction (pre-processing)' **Problem:** The evidence supports a narrower statement than 'analytically inert' implies. It was checked on one model's top-1 family on the prompt side only — not propagated through the four similarity metrics, the valence-sign denominator, or the emotional/non-emotional split. And it is materially untrue for two records: 750 and 751 were reduced from two different people's full names to the bare token '[name]', which then enter the non-emotional control as records with no analysable content. 'Zero top-1 families changed' holds for them only because both were already neutral. **Use instead:** "Redaction changes no top-1 emotion family and moves no GoEmotions probability by more than 0.012, so the prompt-side labels are unaffected. Two records (750, 751) consisted of a name alone and reduce to the token '[name]'; they carry no analysable content and should be excluded rather than counted in the non-emotional group. The effect of redaction on the downstream similarity metrics was not separately verified."

18. "Redaction is also a stronger position than withholding records. A hidden set is itself recoverable, because anyone can probe the public integer ids a

Where: project_dossier.tex 'Name redaction', closing paragraph **Problem:** The argument does not hold. Probing ids reveals which records are absent; it does not recover their contents, which is the disclosure that matters. Withholding eight records leaks eight id numbers; publishing them redacted leaks eight redacted prompts plus their drawings. Either may be defensible, but the stated reason is a non-sequitur dressed as a security argument, and any examiner with an information-security background will notice. The genuine justification — supervisor direction plus analytic completeness — is already available in the document. **Use instead:** "Records are redacted rather than withheld, following supervisor direction, so that the analysed corpus and the demonstrator show the same 1,012 records and cannot diverge; withholding would also have removed those records from the analysis. Redaction does not eliminate re-identification risk entirely, since the drawing and the residual prompt text remain."

19. "An automated test suite of 28 checks covers the environment, data integrity, every analysis, a live run of the pipeline on fresh images, the redactio

Where: project_dossier.tex 'Reproducibility'; PRESENTATION_TECHNICAL_STORY.md '28 automated tests, all passing' **Problem:** The count is wrong (test.sh contains 27 `check` invocations plus two other pass-registering lines = 29 assertions), and 'covers every analysis' overstates what the assertions do. Most are of the form `grep -q 'face images' /tmp/af.log` or `grep -q 'CONSISTENCY' /tmp/ca.log` — they assert that a heading was printed, not that a number is correct. Not one assertion checks a numeric result against an expected value. Combined with the admitted fact that make_figures.py hardcodes its numbers, the suite would pass unchanged if every reported statistic silently changed. **Use instead:** "The automated self-test has two layers: shell smoke checks covering environment setup, dataset and record counts, the redaction audit, a live three-image pipeline run and successful execution of each analysis script; and `test_numeric.py`, which pins every reported statistic to the result file that produces it and unit-tests the metric functions. `make_figures.py` reads the same result files, so a change in the analysis fails the build." (Superseded 9 September: the smoke-only state described in the Problem above no longer applies. Do not quote an assertion count.)

20. "Independently reviewed for statistical validity; over-claims corrected."

Where: emoviz/outputs/RESULTS_final.md, line 7 (header block) **Problem:** No reviewer, date, scope or record of the review is given anywhere in the project. In a dissertation this is an unverifiable appeal to authority occupying the position where a method statement should be, and it invites the examiner to ask who reviewed it — a question with no good answer if the reviewer was an LLM or the student. It also sits awkwardly beside the demonstrable statistical problems above (a degenerate null reported as a validity check, a headline with no null, no test of the headline difference), which suggests the review was not exhaustive. **Use instead:** Delete the sentence, or replace it with a factual audit trail: "Results were re-audited on 1 September 2026; three corrections were made and are recorded below — the positive-versus-negative direction was withdrawn, a label-space mismatch that deflated the baseline's top-3 overlap was fixed (15%->23%), and the face-detection rule was changed to token matching."

21. "A negative control — ~400 prompts contain no feeling … The method scores exactly chance on them"; "Every number is shown against chance."

Where: PRESENTATION_TECHNICAL_STORY.md sections 6 and 9; emoviz/README.md 'Interactive interface' **Problem:** 'Exactly chance' overstates a null result: z=+0.3 means indistinguishable from chance at this sample size, not equal to it, and for two of the four metrics it is an identity rather than a measurement. 'Every number is shown against chance' is contradicted by the project's own consistency headline (68.8%), which appears with no null in four documents. Since the presentation explicitly nominates sections 6 and 7 as the parts that 'read as research', an examiner finding the one un-nulled headline sitting between them is costly. **Use instead:** "A negative control: on ~400 prompts the text reader labels neutral, valence-arousal agreement is 76.6% against a chance level of 76.6% (z=+0.3) — indistinguishable from chance. Every reported number is quoted against its measured chance level, including the consistency result (68.8% against a 59.4% null, z +4.5)."

22. "BLIP-VQA positivity bias: reads ~66% of faces as smiling"

Where: RESULTS_final.md 'Robustness & honesty notes'; RESULTS_describe_extract.md 'Honest limitations' ("66% of all faces as smiling/happy (676/1,025 expression answers)") **Problem:** Wrong denominator, and the parenthesis gives it away: 676/1,025 is the share of ALL records, including the 421 in which no person was reported. Among the 591 records classified as faces the rate is 57.4%. The dossier and the presentation use the correct 57% figure, so the project states two different values for one statistic in different documents — with the inflated one in the file labelled authoritative. **Use instead:** "BLIP-VQA positivity bias: it reports a smiling or happy expression for 57% of the 591 records classified as containing a readable face (337/591), and for 66% of all 1,025 expression answers — including 332 of the 421 records where it simultaneously reported no people present, which is itself evidence of hallucination on this question."

---

3. Inconsistencies, and the value to standardise on

Quantity	Correct value
BLIP-VQA smiling / positivity rate	57% (337/591, token match on smile/smiling/happy) of face-classified records; 66% of all expression answers (669/1,012 analysed; 676/1,025) regardless of whether a person was reported; 332 of the 414 no-people scene records. Standardise on "57% of face-classified records" and always state the denominator.
Number of literature themes	Eight themes. The 97 read / 24 excluded / 73 kept figures are correct and match FILTERING_LOG.md.
Size of the automated test suite	Do not quote a count. Since 9 September the suite has a numeric layer (`test_numeric.py`) that pins every headline value to its result file; describe it as "the automated self-test".
Whether all headlines survive the distinct-prompt robustness check	On 763 distinct prompts (seed 7, 300 shuffles): all-emotional z +9.3 to +10.8; face valence-sign 71.8% (z +8.4); scene valence-sign 37.3% (z +3.5); scene family 30.8% (z +2.6, which does not survive Holm–Bonferroni). The scene headline does not survive at +9 to +11. Delete the blanket sentence.
Distinct-prompt face z-score	71.8%, z +8.4 (seed 7, 300 shuffles — the authoritative run in `analysis_final.py` / `analysis_final.json`). Quote that one value.
Emotion-family agreement on emotional prompts reported as two different numbers	Both are correct but they are different measurements: 53.0% (322/607 — not 53.1) uses the GoEmotions prompt reader; 49.3% uses the BART-MNLI prompt reader (similarity_v2.py reads r['prompt_emotions']). Superseded 9 September by the like-for-like `method_comparison.py`: 53.0% vs 22.4%, +30.6 pp [+25.5, +35.6], cluster-level $p < 0.0001$. Label the second explicitly "NLI prompt side, baseline-comparison subset".
"Top-3 soft overlap" is computed by two different formulas under one name	Two distinct metrics. Use 50.6% as the top-3 overlap for the main analysis and either rename or recompute the baseline comparison with the same formula; as it stands 48.6% vs 22.8% is not on the same scale as the headline.

Quantity	Correct value
Definition of the 'emotional prompt' subset used in the baseline comparison	The 608 records are selected by the v1 lexicon rule in emoviz/text_emotion.py (LexiconBackend.is_emotional), not by the GoEmotions rule that defines the 607. The near-identical counts conceal two different definitions of the same construct. State the v1 rule, or re-run the comparison on the GoEmotions-defined 607.
GoEmotions label-space size	28 labels (27 emotions + neutral) mapped to 7 families. The code's G2F dictionary has exactly 28 keys. Use "28 -> 7" everywhere.
Number of Ekman families in the similarity ladder	Seven (Ekman's six plus neutral). PIPELINE_DESIGN.md carries an AS-BUILT banner but the '-6' still appears and will be read.
Per-record pipeline runtime	~1.5 s/record nominal, degrading to ~12 s/record under memory pressure (~30-90 min total). Quote the range — the degradation is part of the engineering story.
Repeated-prompt group definition and counts	87 distinct prompts drawn 2+ times, covering 336 records, of which 249 are repeats of an earlier record; largest group 'happy' x42. The 65-run definition is superseded. The presentation should read "249 records repeat an earlier prompt", not "249 prompts".
Share of the corpus written entirely in lower case	77.6% of all 1,025 records; 78.5% of the 1,012 analysed. Say "about 78% of the analysed records" — every other count in the documents is on the 1,012 basis.
Neutral-labelled control z-score	z +0.3 (seed 7) for the valence-arousal metric only; +0.04 on the cleaned 335 after removing 68 affect-word prompts and 2 name-only records (`audits.py`). For valence-sign and family the z is identically 0 by construction (degenerate null), not by measurement. "z ~ 0 on every metric" implies four independent passes when there is one.

4. Terminology — one term per concept

Nine terms are used interchangeably for things that are not the same, and the drift is not cosmetic — in several places it is what produces the over-claim.

1. ALIGNMENT / AGREEMENT / SIMILARITY / FIDELITY / MATCHING / CONVEY / CARRY. All seven currently appear for the central quantity ("alignment" 4x dossier, 6x RESULTS_final, 7x RESULTS_describe_extract; "agreement" 8x/4x/4x; "similarity" 7x/2x/8x; "fidelity" in both results files; "convey" 4x dossier, 2x README). STANDARDISE ON: AGREEMENT, always with both readers named — "prompt-reader/image-reader agreement". Why: it is the only one of the seven literally true of what was computed, two automatic labellers landing on the same label. "Alignment" is borrowed from text-to-image evaluation, where it means correspondence to the prompt's actual content — the thing this project cannot measure. "Fidelity" asserts faithfulness to an original. "Convey" and "carry" assert perception by a viewer, the project's single most important disclaimed limitation, and yet they appear in a figure title. Keep SIMILARITY for one narrow use only — the valence-arousal metric, which really is a graded distance — never as a synonym for the whole ladder; the supervisor's "similarity not matching" framing then correctly describes the metric design, which is what she meant. Retire "fidelity" and "matching" except when saying what the project does NOT do.

2. MOOD. PRESENTATION_TECHNICAL_STORY.md introduces a ninth name — "74% mood agreement", "mood-direction" — for what every other document calls valence-sign agreement. STANDARDISE ON: VALENCE-SIGN AGREEMENT, glossed once as "do both readers land on the same side of pleasant/unpleasant".
3. FAMILY / CATEGORY / CLASS / LABEL. The dossier uses "category" for GoEmotions' 27 and "family" for the collapsed 7; PIPELINE_DESIGN says "Ekman-6 family"; the code says "fam". STANDARDISE ON: CATEGORY for the 28 GoEmotions output labels, FAMILY for the 7 collapsed targets, always written "seven emotion families", never "Ekman-6". Why: the two levels are different objects and the reader must know which one agreement is being measured on.
4. FACE / SCENE. These labels claim more than the rule delivers, and the mislabelling is what produced the wrong circularity explanation. STANDARDISE ON: FACE-CLASSIFIED and SCENE-CLASSIFIED — better still, PERSON-REPORTED and NO-PERSON-REPORTED, since 414 of 421 "scenes" are records where the describer answered "none" to the people question. Why: under the current names, "faces carry emotion" reads as a claim about faces in the world; under the corrected names it reads as a claim about a classifier's output, which is what it is.
5. NULL / CHANCE / BASELINE. "Null", "null baseline", "chance line", "chance level" and "baseline" all appear — and "baseline" additionally means the frozen direct-CLIP method. STANDARDISE ON: RANDOM-PAIRING NULL for the permutation reference (or "chance level" in the presentation, defined once), reserving BASELINE exclusively for the direct-CLIP comparison method. Why: "baseline comparison" currently means two different things two paragraphs apart in RESULTS_final.md.
6. DESCRIBE->EXTRACT / describe-then-extract / "the new method" / "Chen's architecture" / "the supervisor's architecture". STANDARDISE ON: DESCRIBE-THEN-EXTRACT, used every time, with the attribution made once in the methodology chapter. Why: "the supervisor's architecture wins" as a findings heading reads as deference rather than evidence, and the examiner marks the student, not the supervisor.
7. RECORD / IMAGE / DRAWING / PROMPT / ENTRY / SELF-REPORT / "text-image pair". All seven appear as the unit of analysis, and this is exactly what lets "1,025 images" slide into "about a thousand members of the public". STANDARDISE ON: RECORD as the unit (a prompt-drawing pair), DRAWING for the image, PROMPT for the text. Never pair a person-word ("visitors", "members of the public") with a record count.
8. RUN / REPEAT / DUPLICATE / "repeated prompt". data_audit_consistency.py says "runs", RESULTS_final says "repeated prompts", PIPELINE_DESIGN says "duplicate-prompt runs", and a superseded 65-run definition is still in circulation. STANDARDISE ON: REPEAT GROUP (87 of them, 336 records), reserving "duplicate" for the deduplication robustness check only.
9. SIGNIFICANT / z-SCORE. "Significant", "highly significant" and "z" are used as though a permutation z were a p-value, with an explicit "significance threshold (z~2)" drawn in Figure 3 and no correction across roughly forty reported z values. STANDARDISE ON the explicit form: "exceeds its random-pairing null by X standard deviations of the permutation distribution (permutation $p < .001$, uncorrected)", plus one sentence stating that no multiple-comparison correction has been applied. Why: "significant" without a stated test, direction and correction is the phrase most likely to be circled in red.

---

5. Supervisory direction — what is satisfied, and what is not

Directive	Status
18 Jun — "Get familiar with the previous datasets"	**satisfied**
18 Jun — "Start literature review (e.g., text-to-image)"	**satisfied**
18 Jun — "Get the template"	**partly satisfied**
9 Jul — "Emotions extracted from text analysis; image analysis" (the two strands)	**satisfied**

Directive	Status
9 Jul handwritten — “Cross verify: Blackboard -> Dissertation -> Refer Blackboard” (confirm the requirements and the 16 Sep deadline against Blackboard)	**partly satisfied**
9 Jul handwritten — “6-Section:- Abstract (later) - Introduction: Explain the context - Literature Review: Topic -> Identify the Research Gap - Methodology - Res	**not satisfied**
9 Jul handwritten — “Each Report: 1. How much - 2. Quality - 3. Writing” (what the report is marked on)	**partly satisfied**
9 Jul handwritten — “Own Opinion Doesn’t Matter -> Cite paper -> Evidence Based”	**partly satisfied**
9 Jul handwritten — “Literature Review: Good Structure - Backing Up - Justifying the Result”	**partly satisfied**
9 Jul handwritten — “Implementation: Method & Result later. Literature Review + Intro -> write in Parallel”	**partly satisfied**
9 Jul handwritten — “Start with Review Papers -> Good Review Article”	**partly satisfied**
9 Jul handwritten — “Writing: Don’t push till last minute”	**not satisfied**
30 Jul #1 — SIMILARITY, not matching; group emotions into families; emotion is subjective so binary matching is too simplistic	**satisfied**
30 Jul #1b — “HIGHLIGHT THIS AS THE NOVEL CONTRIBUTION”: others do one-to-one match/accuracy, you do flexible similarity	**partly satisfied**
30 Jul #2 — do not force CLIP to output emotion; DESCRIBE the image, then extract emotion from the description with a separate text model, then compare text-to-	**satisfied**
30 Jul #2b — take the top-3 emotions each side; hallucination is OK; compare the sets for similarity, not exact overlap	**satisfied**
30 Jul #3 — explore drawing/sketch-specific models (trained on drawings/clip-art, shape-focused) to read objects in line art; consider combining CLIP + a sketch	**not satisfied**
30 Jul #4 — faces vs scenes: face images are easier for emotion, scene-only harder; handle the split	**satisfied**
30 Jul #5 — group the dataset: separate emotional-text prompts from non-emotional (object/thing) prompts and analyse them separately	**satisfied**
30 Jul #6 — ethics/pre-processing: KEEP illness and grief content (not identifying on its own, essential to the analysis); REDACT real names; document the proce	**satisfied**
30 Jul #7 — build in a validity evaluation: an ablation (remove a component / halve the data), a comparison to another method, or human feedback	**satisfied**
30 Jul #7b / 17 Jul Q7 / 20 Aug — the human-rating validation study	**superseded**

Directive	Status
30 Jul #8 — start the Literature Review and Introduction NOW, bit by bit; write your OWN draft, use AI ONLY to polish/proofread (AI first drafts are obvious to	**not satisfied**
30 Jul #9 — by 4 Sep the main RESULTS are done: methods, results, good figures, a clear story, a stated contribution; "pretty much done", not fully written	**satisfied**
30 Jul #9b / #10 — fortnightly supervision, weekly if needed in late August; supervisor engagement is itself a marked criterion	**partly satisfied**
Standing open item, asked at 17 Jul (Q1) and again at 20 Aug — was the Festival of the Mind dataset collected with consent covering secondary research use?	**not satisfied**
17 Jul Q3 — what was the full prompt actually sent to DALL-E (the appended style instruction), and which DALL-E version?	**not satisfied**
17 Jul Q6 — are the duplicate-prompt runs independent batch calls or one person hitting regenerate? (affects independence of every per-record statistic)	**partly satisfied**

Outstanding items that would cost marks if left

- **9 Jul handwritten — "6-Section:- Abstract (later) - Introduction: Explain the context - Literature Review: Topic -> Identify the Research Gap - Methodology - Result (short**
  No dissertation document uses this structure. report/chapters/ still holds the template's five Lorem-ipsum chapters (chapter1-chapter5 = Introduction, Literature Survey, Analysis, Planning, Conclusions). The nearest thing built — PROJECT_EXPLAINED.pdf / emoviz/outputs/report_src/project_dossier.tex — runs question -> why the question is open -> dataset -> method -> results -> limits, which is compatibl
- **9 Jul handwritten — "Writing: Don't push till last minute"**
  Today is 3 Sep 2026 — 1 day to the presentation, 13 days to submission — and the dissertation is unwritten: report/chapters/* are Lorem ipsum, last modified 30 Jul 2026. Her 30 Jul restatement (#8, "start Lit Review + Intro NOW, bit by bit") was also not acted on. This is the project's principal risk and it is a direct, twice-repeated instruction.
- **30 Jul #3 — explore drawing/sketch-specific models (trained on drawings/clip-art, shape-focused) to read objects in line art; consider combining CLIP + a sketch-object mo**
  Designed but never run. emoviz/PIPELINE_DESIGN.md §"Stage 1b — sketch/drawing-specific models (Chen's 'shape-focus' idea)" lists kmewhort/beit-sketch-classifier, QuickDraw ViT finetunes and Sketch_LVM/Sketch-CLIP with a pre-written "honest finding to report", but a grep of every .py in the project finds no sketch model anywhere in the code. The rejected-approaches table (PRESENTATION_TECHNICAL_STO
- **30 Jul #8 — start the Literature Review and Introduction NOW, bit by bit; write your OWN draft, use AI ONLY to polish/proofread (AI first drafts are obvious to markers an**
  Not started. report/chapters/chapter1.tex (Introduction) and chapter2.tex (Literature Survey) remain Lorem ipsum, last modified 30 Jul 2026 — the day she gave the instruction. The AI-use rule is also the binding constraint on how the remaining 13 days are used: the raw material now exists in abundance (RESEARCH_GAP.md, ANNOTATED_BIBLIOGRAPHY.md, RESULTS_final.md, PROJECT_EXPLAINED.pdf, the technic
- **Standing open item, asked at 17 Jul (Q1) and again at 20 Aug — was the Festival of the Mind dataset collected with consent covering secondary research use?**

Never answered. ETHICS_APPLICATION_DRAFT.md B2 still reads "[CONFIRM WITH SUPERVISOR — the data were collected at a public event; state the consent basis for research reuse here.]", and the checklist item " The one thing to confirm with Chen" is unticked. Meanwhile the presentation now asserts a position on her behalf — slide 16 speaker note: "No participants and no new data, so no ethics approv

- **17 Jul Q3 — what was the full prompt actually sent to DALL-E (the appended style instruction), and which DALL-E version?**
  Never resolved. Every image is monochrome line art, so a style instruction was appended by the installation software and is absent from data.json; the C2PA manifests say only "DALL-E", not "DALL-E 3", with no seed and no rewritten prompt, so exact regeneration is impossible (memory/dissertation-dataset-quirks.md; MEETING_BRIEF_17Jul2026.md §1). It is carried as a limitation rather than an answer —

---

6. Chapter structure to use

The supervisor specified a six-section structure in her handwritten notes of 9 July, which differs from the template's five chapters. Use hers:

1. **Abstract** — written last
2. **Introduction** — explain the context
3. **Literature Review** — survey the topic, then identify the research gap
4. **Methodology**
5. **Results** — she noted this should be the shorter chapter
6. **Discussion**

Her stated marking emphasis, verbatim: *"Each report: 1. How much 2. Quality 3. Writing"*, and *"Own opinion doesn't matter -> cite paper -> evidence based"*. For the literature review specifically: *"good structure, backing up, justifying the result"*.